

Introduction to \mathcal{R}

Session 2. Data Management

Dag Tanneberg¹

Potsdam Center for Quantitative Research
University of Potsdam, Germany
November 8/9, 2019

¹Quantitative Social Scientist www.dtanneberg.de

Introduction

Words of Warning

Learning to manage data in \mathcal{R} is hard.
Everybody struggles. Reasons include:

- Where's my spreadsheet?
- Barrage of new concepts
- Explosion of functions and notation
- "Too" many options

Stay with me, and...



Outline

- 1 Introduction
- 2 Data Structures
- 3 R Notation
- 4 A Practical Data Management Challenge

Data Structures

Organizing Principles²

■ Dimensionality

How many qualities allowed?

■ Homogeneity

All data single-typed?

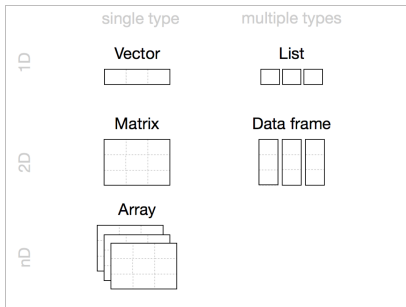
■ Attributes

What metadata available?

■ Today's Focus

Vectors, data frames, and attributes

Figure 1: Typical Data Structures



²Grolemund, G. 2014. Hands-on Programming with R. Sebastopol: O'Reilly, 62.

Atomic vectors: Creation, Dimensionality, Homogeneity

■ Creation

```
die <- 1:6 # Remember our first session?  
die <- c(1, 2, 3, 4, 5, 6) # Standard assignment: c()
```

■ Dimensionality 1D set of data points

■ Homogeneity All data must have the same type.

- Logical: TRUE or FALSE values
- Integer: Whole numbers (\mathbb{Z})
- Double a.k.a. Numeric: Real numbers (\mathbb{R})
- Character: Text data

```
logic <- c(TRUE, FALSE, FALSE); typeof(logic)  
int <- c(-1L, 99L); typeof(int)  
dbl <- sqrt(8L); typeof(dbl)  
chr <- c(TRUE, "Hello", 2.0, "R"); typeof(chr)
```

Atomic Vectors: Attributes

- Provide information *about* your vector
- Atomic Vectors have attributes type, length, and names.
 - type: What kind of data is it?
 - length: How many elements does your vector contain?
 - names: What do you call an element?

```
die <- 1:6
length(die) # Check the length of an atomic vector
length(c(die, die)) # Will the length differ?
names(die) # Explain the output.
names(die) <- c(
  "one", "two", "three", "four", "five", "six"
)
die; die + 1
names(die) <- NULL # Reset the attribute
```


Data Frames: Creation

■ Creation

```
students <- data.frame(  
  first_name = c("Alex", "Jessy", "Barbara", "Jacob"),  
  student_id = c(349857, 796245, 143577, 987456),  
  passed = c(TRUE, TRUE, FALSE, FALSE),  
  stringsAsFactors = FALSE # Will be explained later  
); students
```

##	first_name	student_id	passed
## 1	Alex	349857	TRUE
## 2	Jessy	796245	TRUE
## 3	Barbara	143577	FALSE
## 4	Jacob	987456	FALSE

Data Frames: Dimensionality, Homogeneity, Attributes

■ Dimensionality

- 2D table of grouped vectors
- Vectors must have equal length (or will be recycled)

■ Homogeneity: Data can be of any type.

■ Attributes: (at least) names and dimensions

- Names: What are your rows and columns called?
- Dimensions: No. of rows and columns

Brief description

```
str(students)
```

Dimensions

```
dim(students); nrow(students); ncol(students)
```

Names

```
rownames(students); names(students) # also: colnames()
```

Data I/O

■ Input

- **Easy way:** RStudio's import wizard
- Structured way: `read.table()` and variants or `package:foreign`

■ Output

- **Easy way:** Just kidding. There is no easy way.
- Structured way: `write.table()`, `save()`, or `package:foreign`

```
write.csv(  
  x = student_data,  
  file = "/PATH/TO/DATA/student_data.csv"  
)
```

■ Learn more

- The 'R Data Import/Export' manual (see Help in RStudio)
- <https://www.statmethods.net/input/index.html>

Break

Let's catch our breath and take a 5 minute break.

R Notation

So far, we have seen. . .

- what data structures exist,
- what their elementary properties are,
- what data structures are most important.

The Big Question

How do we access individual values inside our data?

- **Answer:** Indexing
- \mathcal{R} provides numerous ways to index data.
- Most boil down to square brackets: `data[]`.
- Dimensionality defines the number of indexes required
 - Vectors are 1D. \rightarrow `data[i]`
 - Data frames are 2D. \rightarrow `data[i, j]`

Integers

■ *Positive Integers:*

- Example: `data[i, j]`, e.g. `data[1, 2]`
- returns the data indexed by `i` and `j`

```
student_data[ , ] # Select the 1st row & col  
student_data[ , ] # Select rows 1-5 & cols 1-3  
student_data[0, 0] # Explain the result
```

■ *Negative Integers:*

- Example: `data[-i, -j]`, e.g. `data[-1, -2]`
- returns all data **but** `i` and `j`

```
student_data[ , ] # Deselect the 1st row & col  
student_data[ , ] # Select rows 1-5, drop cols 2 & 5  
student_data[-3:4, 1:2] # Explain the error.
```

Blanks

- Use a blank space to extract every value in a dimension
- Example: `data[i,]`

Exercise: Fill in the blanks

`student_data[,]` *# Select the 1st row & all cols*

`student_data[,]` *# Select all rows for cols 2 & 5*

`student_data[,]` *# Return the entire data frame*

Logical Values

- Supply a vector of TRUE & FALSE as your index
- Example: `data[c(TRUE, TRUE, FALSE),]`

```
student_data[FALSE , ] # Explain the result.  
student_data[ , ] # Return only cols 1, 2 & 6  
student_data[student_data[, 2] == 1, ] # Explain.
```

Names

- Supply character vectors to select on names
- Requires attribute names
- Example: `data[, "variable"]`

```
student_data[ , ] # Return cols student & student_id  
student_data[ , "Student"] # Explain the error.  
student_data[ , -"student_id"] # Explain the error.
```

Dollar Signs & Double Brackets

- \$ and [[]] define a second indexing system
- Applies only to lists and data frames
- Return atomic vectors
- \$ requires names, but [[]] always works
- Examples: data\$variable; data[[variable]]; data[[1]]

```
student_data$ # Return column arts  
student_data$ # Return column student_id, rows 1:5  
student_data[[ ]] # Extract column gender  
student_data[[c("student", "student_id")]] # Explain.
```

A Practical Data Management Challenge

The Challenge

A group of students has taken exams in Science, Arts, and Literature. We want to combine their scores, grade them, and prepare a list for display on our office door. Along the way we have to meet several obstacles.

Part A

- 1 Some grades are missing. Angela scored 603 in Science, and Cheryl 28 in Literature. Find and fill the gaps.
- 2 Joel withdrew from class. Remove him from the dataset.
- 3 Gender was coded accidentally. Drop the variable.
- 4 The grades are on widely different scales. **Scale** them to comparable units.

Part B

- 1 Take the average of scaled science, arts, and literature grades for each student.
- 2 Find its .2, .4, .6, and .8 quantiles. Assign grades appropriately.
- 3 Greg was caught cheating. He should receive an F.
- 4 Create a new data frame `grade_data`. This data frame should:
 - include only `student_id` and the final grade;
 - be ordered by `student_id`.

Congratulations

You just learned a hell of a lot of \mathcal{R} . Enjoy lunch!